

How Well-Calibrated Are Driving World Models? A Benchmark for Predictive Uncertainty

Anonymous WACV Algorithms Track submission

Paper ID *****

Abstract

World models for autonomous driving are increasingly used for planning and simulation, yet their predictive uncertainty remains poorly understood. While these models can generate realistic future states, overconfident or poorly calibrated predictions can lead to catastrophic downstream planning failures. To address this gap, we introduce CalibraDrive, the first systematic benchmark for evaluating the calibration of driving world models. We evaluate representative architectures, including 3D occupancy (OccWorld) and video generation (Vista) models, across spatial dimensions, temporal horizons, and scenario difficulties on the nuScenes dataset. Our findings indicate that all tested models are significantly overconfident ($ECE > [TODO: 0.05]$), with calibration degrading further over longer prediction horizons and at greater distances from the ego vehicle. Furthermore, we demonstrate that simple post-hoc recalibration methods can reduce the ECE by over $[TODO: 50\%]$. Importantly, integrating calibrated uncertainty measurably improves downstream planning safety, highlighting the critical need for uncertainty-aware world modeling in autonomous driving.

1. Introduction

World models for autonomous driving are transforming the landscape of vehicle planning and simulation [3, 5, 11, 13, 16]. By learning a predictive distribution over future driving states conditioned on current context and ego actions, these generative models offer an elegant, scalable alternative to hand-crafted simulation environments. However, while world models have demonstrated impressive capability in generating realistic multi-modal futures, a critical question remains largely unanswered: do we know when these models are wrong?

When a driving world model makes an overconfident but incorrect prediction about the future trajectory of a pedestrian or the occupancy of an intersection, downstream planning algorithms that rely on these predictions can make unsafe, potentially catastrophic decisions. Predictive uncertainty is essential for safe autonomous driving [7]. In practice, an autonomous vehicle should act cautiously when its predictive model is uncertain, prioritizing safety over progress. But if the world model itself produces poorly calibrated probabilities—predicting high confidence for incorrect futures—the planner is effectively flying blind.

Calibration has been extensively studied in standard deep learning tasks like image classification and object detection [4], yet it remains largely unexplored for generative driving world models. Existing literature on uncertainty in autonomous driving primarily focuses on perception-level uncertainty or traditional prediction models, lacking a systematic evaluation of modern large-scale world models. While recent methods like C3 [7] and RELIOCC have proposed specific calibration techniques for certain prediction tasks, there currently exists no standardized benchmark for evaluating and comparing the calibration of diverse driving world model architectures.

In this paper, we present CalibraDrive, the first systematic benchmark for evaluating the predictive uncertainty and calibration of driving world models. Our framework spans diverse representative architectures, evaluating both 3D occupancy forecasting (OccWorld [16]) and video generation (Vista [3]) paradigms. Through extensive spatial-temporal analysis, we reveal where and when these models fail, demonstrating that confidence is poorly aligned with accuracy, particularly at longer prediction horizons and in hard driving scenarios.

In summary, our key contributions are:

1. CalibraDrive Benchmark: We introduce the first comprehensive framework for evaluating calibration

[TODO: Include Teaser Figure]

Figure 1. The CalibraDrive benchmark framework. We evaluate the calibration of driving world models (e.g., occupancy and video generation) across spatial and temporal dimensions. Poorly calibrated models lead to unsafe planning, while calibrated uncertainty significantly improves downstream safety.

in generative driving world models across different architectures (occupancy and video).

2. Spatial-Temporal Calibration Analysis: We provide an in-depth analysis of where and when driving world models are miscalibrated, showing degradation over prediction horizons and spatial distances.
3. Recalibration and Downstream Safety: We demonstrate that simple post-hoc recalibration techniques not only significantly improve calibration metrics (e.g., reducing Expected Calibration Error) but also measurably enhance the safety of downstream planning algorithms.

2. Related Work

Driving World Models. The success of large-scale generative models has spurred significant interest in developing world models for autonomous driving. These models generally predict future states conditioned on historical context and ego-vehicle actions, enabling both closed-loop simulation and model-based planning. OccWorld [16] predicts future 3D occupancy and ego trajectories autoregressively. Video-based world models such as Vista [3], GAIA-1 [5], and DriveDreamer [11] formulate the problem as future video generation. Other notable works include UniSim [13], GenAD [15], Copilot4D [14], and Drive-WM [12]. While these models achieve impressive fidelity in predicting future states, their predictive confidence and uncertainty have not been rigorously evaluated.

Calibration in Deep Learning. Calibration measures the alignment between a model’s predicted confidence and its actual empirical accuracy. Deep neural networks, particularly large modern architectures, are often notoriously overconfident [4]. Various methods have been proposed to improve calibration, ranging from post-hoc scaling methods like Platt scaling [8] and temperature scaling [4] to distribution-free guarantees using conformal prediction [9, 10]. Other prominent approaches include Bayesian Neural Networks, MC Dropout [2], and Deep Ensembles [6]. While these techniques are well-established for classification and regression, adapting and evaluating them for high-dimensional, generative driving world models remains

an open challenge.

Uncertainty in Autonomous Driving. Uncertainty estimation is critical for safe autonomous driving. While extensive research has addressed uncertainty in 3D object detection, semantic segmentation, and traditional trajectory prediction, exploring uncertainty in comprehensive driving world models is a nascent field. The most closely related work to ours is C3 [7], which explores calibration for conformal trajectory prediction, and RELIOCC, which targets occupancy calibration. However, these works primarily propose specific methodologies for specific tasks. In contrast, CalibraDrive provides a standardized benchmark and systemic evaluation across fundamentally different world model architectures (e.g., 3D occupancy vs. video generation), temporal horizons, and planning settings.

3. Methodology

3.1. Problem Formulation

We define a driving world model W as a probabilistic model that takes historical context c (e.g., past sensor data, maps) and future ego actions a to produce a distribution over future states. Let $Y \in \mathcal{Y}$ be the true future state. The model produces a predictive distribution $P_\theta(Y|c, a)$, from which we can extract a confidence score $\hat{p}$ for a specific prediction $\hat{Y}$.

A model is perfectly calibrated if the confidence $\hat{p}$ matches the empirical probability of the prediction being correct. Formally:

$$\mathbb{P}(Y = \hat{Y} \mid \hat{p} = p) = p, \quad \forall p \in [0, 1]. \quad (1)$$

For 3D occupancy models, we evaluate binary calibration per voxel, where $\hat{p}$ is the predicted probability of the voxel being occupied. For video generation, we assess the predictive variance and probability of semantic classes per pixel.

3.2. Calibration Metrics

To quantify calibration, we employ standard metrics adapted for the spatial-temporal nature of driving predictions.

- **Expected Calibration Error (ECE):** We partition the confidence scores into B bins. The ECE is the weighted average of the absolute difference between accuracy and confidence across bins:

$$\text{ECE} = \sum_{b=1}^B \frac{|B_b|}{N} |\text{acc}(B_b) - \text{conf}(B_b)|, \quad (2)$$

where N is the total number of predictions, and B_b represents the set of predictions falling into bin b .

- Maximum Calibration Error (MCE): The maximum discrepancy across all bins: $\max_b |\text{acc}(B_b) - \text{conf}(B_b)|$.
- Brier Score: A strictly proper scoring rule measuring the mean squared error of probability predictions: $\frac{1}{N} \sum_{i=1}^N (p_i - y_i)^2$.
- AUROC & Reliability Diagrams: We use AUROC to evaluate discriminative ability, independent of calibration. Reliability diagrams visually plot accuracy versus confidence.

3.3. Spatial-Temporal Analysis

Driving predictions are inherently spatial and temporal. A prediction error directly in front of the ego vehicle is significantly more critical than one 100 meters away. Thus, CalibraDrive stratifies evaluation across several axes:

- Spatial Distance: We evaluate calibration in discrete radial distance bins from the ego vehicle: $[0 - 10m, 10 - 30m, 30 - 50m, 50 - 100m]$.
- Temporal Horizon: We analyze calibration across prediction horizons from 0.5s up to 3.0s into the future.
- Scenario Difficulty: We classify scenarios as easy, medium, or hard based on agent density, road curvature, and average time-to-collision (TTC).
- Dynamic vs. Static: We separately evaluate calibration for static background elements versus dynamic traffic participants.

3.4. Post-Hoc Recalibration

We evaluate three post-hoc recalibration techniques to improve the uncertainty estimates of world models without requiring retraining:

1. Temperature Scaling [4]: A single scalar parameter T is learned on a validation set by minimizing the Negative Log-Likelihood (NLL). The calibrated probability is given by $p' = \sigma(z/T)$, where z is the logit and σ is the sigmoid or softmax function.
2. Conformal Prediction [9]: Provides a distribution-free theoretical guarantee that the true outcome falls within a predicted set with a user-specified probability $1 - \alpha$.
3. Histogram Binning: A non-parametric baseline that aligns probabilities within predefined bins to match their empirical frequencies.

3.5. Downstream Planning Evaluation

To understand the practical impact of calibration, we integrate the world model into a Model Predictive Control (MPC) planning framework. The planner minimizes a cost function $J(a)$:

$$J(a) = \mathbb{E}[C(a)] + \lambda \cdot U(a), \quad (3)$$

Table 1. Calibration metrics on nuScenes val.

Model	ECE ↓	MCE ↓	Brier ↓	AUROC ↑
OccWorld	[TODO: 0.XXX]	[TODO: 0.XXX]	[TODO: 0.XXX]	[TODO: 0.XXX]
Vista	[TODO: 0.XXX]	[TODO: 0.XXX]	[TODO: 0.XXX]	[TODO: 0.XXX]

where $C(a)$ represents standard planning costs (e.g., collisions, progress, comfort), and $U(a)$ is an uncertainty penalty derived from the world model’s predictive variance or entropy. We vary the risk-aversion parameter λ from 0 (uncertainty-agnostic baseline) to 5 (highly risk-averse). We evaluate the resulting planner’s collision rate, route progress, and comfort metrics.

4. Experiments

4.1. Experimental Setup

Dataset: We use the standard nuScenes dataset [1], focusing on the validation split. We curate a subset of 1500 diverse scenarios, stratified into easy, medium, and hard difficulties.

Models: We evaluate two representative architectures:

- OccWorld [16]: An autoregressive 3D occupancy world model. We generate stochastic samples via temperature sampling.
- Vista [3]: A diffusion-based video generation world model. We induce stochasticity by varying the initial noise seed.

For each scenario and model, we generate $N = 20$ stochastic samples to estimate predictive uncertainty.

4.2. Overall Calibration (E1)

Our baseline evaluation reveals that both occupancy and video world models are significantly overconfident. As shown in Table ??, OccWorld and Vista yield ECE values of [TODO: 0.12] and [TODO: 0.15], respectively. The reliability diagrams (Figure ??) further illustrate this trend, showing that predictions with $> 90\%$ confidence often exhibit empirical accuracies below [TODO: 70%].

4.3. Spatial-Temporal Analysis (E2-E4)

Spatial Degradation: Calibration error increases dramatically with distance from the ego vehicle. Table ?? demonstrates that for OccWorld, the ECE in the $0 - 10m$ bin is [TODO: 0.05], but degrades to [TODO: 0.18] in the $50 - 100m$ bin.

Temporal Degradation: Figure ?? shows the progression of ECE over the prediction horizon. At 0.5s, models remain relatively well-calibrated, but by 3.0s, accumulated autoregressive errors and diffusion drift cause the ECE to double.

256 Scenario Difficulty and Dynamics: Hard scenarios
257 exhibit significantly worse calibration than easy sce-
258 narios. Furthermore, uncertainty is particularly poorly
259 calibrated for dynamic objects compared to static back-
260 grounds, suggesting models struggle to properly quan-
261 tify uncertainty in complex, highly dynamic interac-
262 tions.

263 4.4. Qualitative Driving Anecdotes & Failure Modes

264 To understand the practical failure modes of driv-
265 ing world models beyond aggregate metrics, we mine
266 and analyze specific real-world driving anecdotes from
267 nuScenes (Figure 2):

268 Case 1: Blind-Zone Pedestrian Emergence (False-
269 Negative Hazard): When a pedestrian steps into a
270 crosswalk from behind a parked high-profile truck, the
271 world model overconfidently predicts free space ($p <$
272 0.15) in the unobserved blind zone instead of express-
273 ing high epistemic uncertainty. For downstream plan-
274 ners, this false certainty suppresses braking, directly
275 inducing safety-critical collision risk.

276 Case 2: Unprotected Left Turn (Mode Collapse):
277 Navigating an intersection with oncoming traffic
278 presents a bifurcating decision boundary. Rather than
279 capturing a bimodal distribution across stochastic roll-
280 outs, the model exhibits mode collapse, predicting with
281 unjustified certainty that oncoming agents will yield.

282 Case 3: High-Speed Lateral Cut-In (Dynamic Lag):
283 When an adjacent vehicle abruptly swerves into the ego
284 buffer, the model’s confidence exhibits a severe tem-
285 poral lag, underestimating lateral displacement during
286 transition timesteps ($t \approx 1.5s$).

287 Case 4: Long-Horizon Hallucination (Autoregressive
288 Drift): At horizons $t \geq 2.5s$, compounding discrete to-
289 ken errors cause phantom obstacles to materialize on
290 clear road surfaces while genuine obstacles fade, ex-
291 plaining the dramatic temporal calibration decay ob-
292 served in Section 4.

293 Front-Camera Behavioral Rollouts: In addition to
294 top-down grid representations, Figure 3 visualizes the
295 driver’s front-camera perspective across sequential roll-
296 outs ($t = 0.5s \rightarrow 3.0s$). Projecting predictive uncer-
297 tainty directly onto the physical road surface and dy-
298 namic actors reveals that world models often exhibit
299 severe epistemic blind spots during dynamic cut-ins,
300 failing to elevate predictive variance until after physi-
301 cal buffer violations occur.

302 4.5. Recalibration (E5)

303 We apply post-hoc recalibration methods and report
304 the results in Table ???. Temperature scaling proves
305 highly effective, reducing the ECE by over [TODO:
306 50%] for both models. Conformal prediction success-

fully provides the desired marginal coverage guaran-
tees (Table ??), yielding reliable uncertainty bounds
for high-stakes predictions.

4.6. Downstream Planning (E6)

To validate the utility of calibrated uncertainty, we
evaluate the MPC planner. Table ?? demonstrates that
integrating calibrated uncertainty ($\lambda > 0$) reduces the
collision rate by [TODO: 30%] compared to the base-
line planner ($\lambda = 0$). Furthermore, Figure ?? illustrates
the safety-progress trade-off: highly risk-averse settings
($\lambda = 5$) eliminate collisions but result in overly conser-
vative driving, highlighting the importance of tuning
the uncertainty penalty based on properly calibrated
confidence estimates.

5. Conclusion

In this paper, we introduced CalibraDrive, the first
benchmark dedicated to evaluating predictive uncer-
tainty and calibration in driving world models. Our
comprehensive analysis of representative occupancy
and video models on the nuScenes dataset reveals a
pervasive issue: modern generative driving models are
significantly overconfident. This miscalibration wors-
ens over longer prediction horizons and in hard driving
scenarios, potentially jeopardizing downstream safety.

Crucially, we show that simple post-hoc recalibra-
tion techniques, such as temperature scaling, effectively
mitigate this overconfidence and are easy to adopt.
Furthermore, integrating these calibrated uncertainty
estimates into downstream planning significantly re-
duces collision rates, highlighting the practical neces-
sity of uncertainty-aware world modeling.

Future work will explore online, dynamic calibra-
tion methods and the integration of calibration objec-
tives directly during the training phase of driving world
models. We hope CalibraDrive will serve as a standard
evaluation axis, encouraging the development of not
just highly realistic, but highly reliable world models
for autonomous driving.

Code and evaluation scripts will be made publicly
available.

References

- [1] Holger Caesar et al. nuscenes: A multimodal dataset
for autonomous driving. In CVPR, 2020. 3
- [2] Yarin Gal and Zoubin Ghahramani. Dropout as a
bayesian approximation: Representing model uncer-
tainty in deep learning. In ICML, 2016. 2
- [3] Shenanyi Gao et al. Vista: A generalizable driving
world model with high fidelity and versatile controlla-
bility. In NeurIPS, 2024. 1, 2, 3

Figure 2. Qualitative real-world driving anecdotes mined by CalibraDrive. We highlight four archetypal corner cases: (Case 1) Blind-Zone Pedestrian Emergence, where models overconfidently predict free space in occlusions; (Case 2) Unprotected Left Turn with oncoming traffic, showing mode collapse; (Case 3) High-Speed Lateral Cut-In, revealing dynamic calibration lag; and (Case 4) Long-Horizon Hallucination at $t = 3.0s$. Columns depict ground truth context, model predicted confidence, predictive uncertainty, and safety-critical failure regions.

Figure 3. Front-camera driver perspective and behavioral dynamics across prediction horizons ($t = 0.5s$ to $3.0s$). Row 1 shows the physical road environment and dynamic actor movements (cut-in vehicle and roadside pedestrian). Row 2 projects the model’s predictive uncertainty heatmap directly onto the camera view. Row 3 depicts synchronized BEV occupancy and ego planned path.

- [4] Chuan Guo et al. On calibration of modern neural networks. In ICML, 2017. [1](#), [2](#), [3](#)
- [5] Anthony Hu et al. Gaia-1: A generative world model for autonomous driving. arXiv, 2023. [1](#), [2](#)
- [6] Balaji Lakshminarayanan et al. Simple and scalable predictive uncertainty estimation using deep ensembles. In NeurIPS, 2017. [2](#)
- [7] Jiyang Mei et al. C3: Conformal calibration for confident trajectory prediction. arXiv preprint arXiv:2512.05927, 2025. [1](#), [2](#)
- [8] John Platt. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. Advances in Large Margin Classifiers, 1999. [2](#)
- [9] Yaniv Romano et al. Conformalized quantile regression. In NeurIPS, 2019. [2](#), [3](#)
- [10] Vladimir Vovk, Alex Gammerman, and Glenn Shafer. Algorithmic learning in a random world. In Springer, 2005. [2](#)
- [11] Xiaofeng Wang et al. Drivedreamer: Towards real-world-driven world models for autonomous driving. arXiv, 2023. [1](#), [2](#)
- [12] Yuqi Wang et al. Drive-wm: Driving in the world model. arXiv, 2024. [2](#)
- [13] Ze Yang et al. Unisim: A neural closed-loop sensor simulator. CVPR, 2023. [1](#), [2](#)
- [14] Yuxuan Zhang et al. Copilot4d: Learning a 4d world model for autonomous driving. In ICLR, 2024. [2](#)
- [15] Wenzhao Zheng et al. Genad: Generative autonomous driving. In ECCV, 2024. [2](#)
- [16] Wenzhao Zheng et al. Occworld: Learning a 3d occupancy world model for autonomous driving. In ECCV, 2024. [1](#), [2](#), [3](#)